



Deep Autoencoder Neural Networks: A Comprehensive Review and New Perspectives

Ibomoiye Domor Mienye¹ · Theo G. Swart¹

Received: 26 June 2024 / Accepted: 28 February 2025 / Published online: 15 March 2025
© The Author(s) 2025

Abstract

Autoencoders have become a fundamental technique in deep learning (DL), significantly enhancing representation learning across various domains, including image processing, anomaly detection, and generative modelling. This paper provides a comprehensive review of autoencoder architectures, from their inception and fundamental concepts to advanced implementations such as adversarial autoencoders, convolutional autoencoders, and variational autoencoders, examining their operational mechanisms, mathematical foundations, typical applications, and their role in generative modelling. The study contributes to the field by synthesizing existing knowledge, discussing recent advancements, new perspectives, and the practical implications of autoencoders in tackling modern machine learning (ML) challenges.

1 Introduction

In today's machine learning era, representation learning is essential for deriving meaningful features automatically from raw data, which is vital for tasks such as pattern recognition and classification [1–3]. As the complexity of tasks in ML grows, from image and speech recognition to natural language processing (NLP), the demand for more refined and adaptive representation learning techniques becomes critical. The quality of these representations plays a vital role in the performance of downstream tasks and directly influences the effectiveness of classification and predictions.

Autoencoders (AEs), a special class of neural networks, are central to the progress in deep representation learning [4]. These networks are designed to compress input data into a condensed set of representations and then reconstruct the input from these representations, ideally with minimal information loss [5]. This dual process of encoding and decoding aids in reducing the dimensionality of data and the learning of valuable features without the need for extensive human supervision. Historically, autoencoders have been vital in pushing the boundaries of unsupervised and semi-supervised

learning, where labelled data is either scarce or costly to acquire [6].

Despite their significance, the literature on autoencoders often remains fragmented, with studies and applications scattered across various contexts and fields. For instance, Vincent et al. [7] utilized denoising autoencoders to demonstrate the potential of autoencoders in feature learning by reconstructing inputs from corrupted versions, thus enhancing the robustness of learned features. In a different application, autoencoders have been employed in anomaly detection within complex datasets, as explored by Sakurada and Yairi [8], where they effectively identified unusual patterns within aerospace telemetry data. Therefore, considering the applications of autoencoders and innovative methodologies, there is a vital need for a comprehensive review that not only traces their development but also examines new methodologies and techniques that have emerged in recent years. Hence, this study aims to bridge that gap and, by doing so, explore both the potential and limitations of autoencoders, providing clearer guidance for future research.

This review is timely given the rapid advancements in deep learning architectures, such as the emergence of variational autoencoders (VAEs) and adversarial training models, leading to a pressing need to reassess the autoencoder domain. Therefore, this review aims to consolidate these diverse applications and innovations, providing a cohesive overview that highlights how autoencoders have evolved and continue to make a significant impact in the deep learning field. The review also offers new perspectives on how

✉ Ibomoiye Domor Mienye
ibomoiyem@uj.ac.za

Theo G. Swart
tgswart@uj.ac.za

¹ Institute for Intelligent Systems, University of Johannesburg, Auckland Park, Johannesburg 2006, South Africa

autoencoders can be effectively utilized in both existing and emerging ML applications.

The rest of the paper is structured as follows: Sect. 2 reviews related works, while Sect. 3 presents an overview of autoencoders, describing the core principles and basic components, such as stochastic gradient descent and backpropagation. Section 4 discusses the evolution of autoencoder architectures, from the basic architectures, such as sparse and denoising autoencoders, to more advanced architectures like variational, adversarial, convolutional autoencoders, and others. Section 5 further discusses variational autoencoders and their role in generative modelling, including the key extensions proposed to enhance their generative abilities. Section 6 presents recent advancements and methodology, including innovations in autoencoder design and emerging training methodology. Meanwhile, Sect. 7 presents new perspectives in autoencoder applications. Section 8 discusses future research directions, while Sect. 9 concludes the paper.

2 Related Works

The study of autoencoders touches on numerous aspects of ML, from basic data compression techniques to more complex applications in unsupervised and generative modelling. Historically, autoencoders have been primarily utilized for dimensionality reduction, similar to principal component analysis (PCA) but with a nonlinear approach [8]. Early studies, such as those by Hinton et al. [9], demonstrated that autoencoders could effectively reduce data dimensionality while preserving its salient characteristics, thus paving the way for their use in more advanced ML tasks.

VAEs and adversarial autoencoders (AAEs) represent two significant advancements in the evolution of autoencoders. Kingma and Welling [10] presented an introduction to VAEs, which provided a new framework for handling the probabilistic aspects of latent variable models, enabling the generation of new data points from complex, learned distributions. This innovation has strengthened the use of autoencoders in generative tasks and also deepened our understanding of the latent space's role in unsupervised learning. Similarly, adversarial autoencoders, introduced by Makhzani et al. [11], combine the autoencoder architecture with the adversarial training methodology of generative adversarial networks (GANs), introduced by Goodfellow et al. [12]. This hybrid approach has enhanced the quality of generated data, making AAEs particularly useful for tasks requiring high-fidelity data generation.

Despite these significant advancements, the research into autoencoders has revealed gaps, particularly in their application in unsupervised and semi-supervised learning frameworks [13]. The integration of autoencoders in these areas remains underexplored, especially in leveraging unlabeled

data for complex predictive modelling. This gap is evident in real-world applications where extensive labelled data is unavailable, such as anomaly detection in network security or in complex disease modelling in healthcare.

Meanwhile, numerous reviews in the literature have explored the broad applications and theoretical foundations of autoencoders. For instance, comprehensive reviews in [14–16] delve into the fundamentals of autoencoder technology, including various architectures and their applications in unsupervised learning. These works provide a solid foundation for understanding the standard autoencoder models but tend to focus primarily on their application within traditional tasks such as image and speech recognition. Additionally, Ng [17] presented an overview of autoencoders, focusing on sparse autoencoders. While foundational, these reviews often do not explore the more recent expansions into generative modelling and the integration of autoencoders with other advanced neural network architectures that are critical in handling more complex, real-world datasets.

In contrast, more recent reviews, such as those in [18, 19], have begun to address the use of autoencoders in novel applications. These reviews expand the discussion into how autoencoders can be adapted for specific industrial applications, highlighting their practical implications. However, they still largely overlook the interdisciplinary potential of autoencoders, particularly in how these models can be tailored to meet specific challenges across different domains such as healthcare, finance, and environmental science.

Therefore, this review is unique as it not only explores the broad spectrum of autoencoder applications from basic dimensionality reduction to advanced generative tasks but also emphasizes their evolving role in cross-disciplinary settings. It explores the integration of autoencoders with other types of neural networks to enhance their predictive accuracy and generative quality. Furthermore, this review contributes extensively to the literature by presenting new perspectives on autoencoder design and applications in sectors that have been less examined in traditional autoencoder literature, such as genomics for personalized medicine and real-time financial anomaly detection.

3 Overview of Autoencoders

Autoencoders are a specific type of neural network designed to learn efficient representations (encodings) of unlabeled data, typically for the purpose of dimensionality reduction or feature learning [20]. By design, autoencoders seek to reproduce their input at the output layer after encoding and decoding it through a hidden layer known as the bottleneck. This bottleneck is a critical feature of the autoencoder architecture, as it forces the network to prioritize which aspects of the input data are most salient, thus learning to ignore

“noise” and less important details [21]. An autoencoder is comprised of two main parts: the encoder and the decoder. The encoder function $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^p$ maps the input $\mathbf{x}$ to a latent space representation $\mathbf{z}$, where typically $p < d$, meaning the latent space has fewer dimensions than the input space. The decoder function $\psi : \mathbb{R}^p \rightarrow \mathbb{R}^d$ then attempts to reconstruct the input, denoted as $\hat{\mathbf{x}}$, from this compressed data, as shown in Fig. 1. The entire autoencoder can be described by the composition of these two functions:

$$\hat{\mathbf{x}} = \psi(\phi(\mathbf{x})) \quad (1)$$

The goal of the autoencoder is to minimize the difference between the input $\mathbf{x}$ and its reconstruction $\hat{\mathbf{x}}$, a criterion that is quantified using a loss function. The most commonly used loss function in autoencoders is the mean squared error (MSE), which is used to quantify the difference between the original and reconstructed inputs [5]. Meanwhile, for binary input data, the cross-entropy loss function is often more appropriate. For a training dataset $\mathbf{X} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}\}$, the MSE and cross-entropy loss functions are defined as:

$$L(\phi, \psi) = \frac{1}{n} \sum_{i=1}^n \|\mathbf{x}^{(i)} - \psi(\phi(\mathbf{x}^{(i)}))\|^2 \quad (2)$$

$$L(\phi, \psi) = -\frac{1}{n} \sum_{i=1}^n [\mathbf{x}^{(i)} \log(\psi(\phi(\mathbf{x}^{(i)}))) + (1 - \mathbf{x}^{(i)}) \log(1 - \psi(\phi(\mathbf{x}^{(i)})))] \quad (3)$$

Meanwhile, regularization terms such as the L1 norm (to encourage sparsity in the latent representation) may also be added to the loss function. Regularization helps to prevent overfitting by penalizing large weights, and when used in autoencoders, it can also help to ensure that the learned representations are sparse, which often improves the interpretability and usefulness of the encoded features [6]. Autoencoders are typically trained using backpropagation and an

optimization algorithm such as stochastic gradient descent. The following subsections discuss these techniques.

3.1 Stochastic Gradient Descent

Stochastic gradient descent (SGD), a widely used optimization algorithm in neural network training, updates the model weights iteratively based on a subset of the entire data set, known as a minibatch [23]. SGD updates the weights $\mathbf{W}$ by moving in the opposite direction of the gradient of the loss function $\nabla L(\mathbf{W})$ with respect to the weights, scaled by a step size η (learning rate). The update rule can be expressed as follows:

$$\mathbf{W}_{t+1} = \mathbf{W}_t - \eta \nabla L(\mathbf{W}_t) \quad (4)$$

where $\mathbf{W}_t$ represents the weights at iteration t , and $\mathbf{W}_{t+1}$ are the weights after the update. The choice of the learning rate η and the batch size are crucial in determining the convergence rate and stability of the training process. A smaller learning rate might lead to slower convergence, while a larger rate could cause oscillations or divergence [24]. Similarly, the batch size affects the estimate of the gradient; larger batches provide a more accurate gradient but require more computational resources [23].

3.2 Backpropagation

Backpropagation is used to compute the gradient of the loss function with respect to each weight in the network, which is essential for updating the weights to minimize the loss [25]. The essence of backpropagation in the context of autoencoders is to apply the chain rule of calculus to compute gradients of the loss function with respect to each weight in the network. For an autoencoder, the gradient of the loss function can be particularly complex due to the composition of the encoder and decoder functions. Given the encoder function ϕ and the decoder function ψ , the gradient of the




Fig. 1 An example of the autoencoder process [22]

loss function L with respect to the encoder weights $\mathbf{W}_\phi$ is given by:

$$\nabla_{\mathbf{W}_\phi} L = \frac{\partial L}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}_\phi} \quad (5)$$

where $\mathbf{z} = \phi(\mathbf{x}; \mathbf{W}_\phi)$ is the encoded representation. For the decoder, the gradient is computed similarly:

$$\nabla_{\mathbf{W}_\psi} L = \frac{\partial L}{\partial \hat{\mathbf{x}}} \frac{\partial \hat{\mathbf{x}}}{\partial \mathbf{W}_\psi} \quad (6)$$

where $\hat{\mathbf{x}} = \psi(\mathbf{z}; \mathbf{W}_\psi)$ is the reconstructed output.

4 Evolution of Autoencoder Architectures

Autoencoders, first introduced in the 1980s, were originally proposed as a data compression algorithm similar to PCA but with the capability to capture non-linear relationships [26]. Hinton and Zemel [27] provided early foundational work on autoencoders and presented how autoencoders could be used for learning data codings in an unsupervised manner. In the 1990s and early 2000s, autoencoders were primarily explored to understand their capacity for non-linear dimensionality reduction [28, 29]. For example, Kramer [30] extended autoencoder application to nonlinear PCA, enhancing their utility in analyzing complex data sets in fields like finance and signal processing.

The potential of autoencoders was further expanded by Hinton and Salakhutdinov [9], who demonstrated their effectiveness in deep architectures for dimensionality reduction. In 2008, Vincent et al. [7] introduced denoising

autoencoders, advancing their capability to remove noise from data, which was especially beneficial for signal and image processing. The development of VAEs by Kingma and Welling [10] marked a significant milestone by introducing a probabilistic framework for generative modelling, allowing for the generation of new, diverse data instances. Finally, Makhzani et al. [11] developed adversarial autoencoders, which integrated adversarial training to further improve the diversity and realism of generated data, proving useful in simulation and training environments.

The fundamental work by Hinton and Salakhutdinov [9] was critical in demonstrating that deep neural networks could effectively reduce dimensionality by learning to encode inputs into a lower-dimensional space and then reconstructing the original input from this compressed representation. The key timelines in autoencoder evolution are shown in Fig. 2. Since these key developments, autoencoders have seen diverse applications across numerous fields, including image processing, natural language processing, and anomaly detection [31, 32]. However, researchers have continued to innovate on the architecture and applications of autoencoders, pushing the boundaries of what can be achieved with these models.

4.1 Basic Architectures

4.1.1 Sparse Autoencoders

Sparse autoencoders (SAE) are a variant of autoencoders that add a sparsity constraint to the hidden layers, which encourages the network to activate only a small number of neurons at any given time [20]. This is designed to improve




Fig. 2 Key Timelines in Autoencoder Evolution

the ability of the autoencoder to capture more meaningful and descriptive features in its hidden representations. The introduction of sparsity into autoencoders was significantly influenced by Olshausen and Field [33], which focused on natural image statistics. They demonstrated that sparse representations could lead to more efficient and robust feature detection in neural networks, mirroring the response properties of neurons in the mammalian visual cortex. Their findings provided a strong theoretical foundation for using sparsity as a tool to improve the representational power of autoencoders.

Practically, sparsity in autoencoders is typically enforced by adding a sparsity penalty to the loss function during training. This penalty encourages the network to reduce the activity of its neurons [17]. One common approach to implementing this penalty is to use the Kullback–Leibler (KL) divergence to measure the discrepancy between the actual activation of the neurons and a desired sparsity level ρ . The overall loss function for a sparse autoencoder can then be expressed as:

$$L_{\text{sparse}} = L + \beta \sum_k \text{KL}(\rho \| \hat{\rho}_k) \quad (7)$$

where L is the reconstruction loss, β is a coefficient determining the weight of the sparsity penalty in the total loss function, ρ is the sparsity target (a small value close to zero), and $\hat{\rho}_k$ is the average activation of the k -th neuron in the hidden layer [17]. The KL-divergence is calculated for each neuron to ensure that the average activation remains close to the target sparsity level, pushing most neurons to remain inactive most of the time. Furthermore, incorporating sparsity into autoencoders helps in several ways. Firstly, it leads to the discovery of more statistically significant features in the data as the model learns to represent the entire dataset with fewer active neurons. Secondly, sparse representations are often more interpretable and can lead to better generalization from training data to unseen data [34, 35]. Finally, by reducing the dimensionality of the data more aggressively, sparse autoencoders can perform better in tasks where the intrinsic dimensionality of the data is low.

4.1.2 Denoising Autoencoders

Denoising autoencoders (DAEs) represent a significant advancement in the field of neural networks, specifically designed to enhance the robustness and utility of learned representations. Introduced by Vincent et al. [7], these models address the challenge of overfitting by forcing the network to focus on the most relevant features of the data. The concept of DAE was developed to counter the tendency of traditional autoencoders to memorize input data rather than to learn to generalize from it. By reconstructing

a clean output from a deliberately corrupted version of the input, DAEs learn to ignore irrelevant stochastic noise in the data. This approach was groundbreaking as it provided a simple yet effective method for improving the feature extraction capabilities of autoencoders without requiring complex modifications to the network architecture. The primary mechanism of DAEs involves the intentional corruption of input data before it is fed into the network. This corruption can take various forms, such as the addition of Gaussian noise, masking some input features, or applying salt-and-pepper noise (i.e., randomly blacking out pixels in images) [7]. The autoencoder then works to predict the clean original inputs from these corrupted data instances during training. This process is depicted in Fig. 3. The general process can be mathematically represented as follows:

$$\mathbf{x}' = q_D(\mathbf{x}) \quad (8)$$

$$\mathbf{y} = \phi_\theta(\mathbf{x}') \quad (9)$$

$$\hat{\mathbf{x}} = \psi_{\theta'}(\mathbf{y}) \quad (10)$$

$$L_H(\mathbf{x}, \hat{\mathbf{x}}) = \|\mathbf{x} - \hat{\mathbf{x}}\|^2 \quad (11)$$

where $\mathbf{x}$ is the original clean input, $\mathbf{x}'$ is the corrupted version of the input, $\mathbf{y}$ is the encoded representation, and $\hat{\mathbf{x}}$ is the reconstruction of the clean input from the corrupted version by the decoder. The function q_D represents the corruption process that stochastically modifies $\mathbf{x}$, ϕ_θ is the encoder function parameterized by θ , and $\psi_{\theta'}$ is the decoder function parameterized by θ' . The training objective is to minimize the loss function L_H , which typically measures the mean squared error between the original and reconstructed inputs, promoting robustness against input noise. DAEs have been effectively utilized in various applications where robustness against input noise is crucial, such as image and audio processing, medical image analysis, and any scenario where data may be incomplete or corrupted by environmental factors [37–39]. The ability to recover data from corrupted inputs makes these models particularly valuable in practical settings where perfect data collection is often unfeasible.




Fig. 3 Basic architecture of the denoising AE [36]

4.2 Advanced Architectures

4.2.1 Variational Autoencoders

Variational autoencoders represent a significant development in the field of neural networks and generative modeling as introduced by Kingma and Welling [10]. Their work provided a robust statistical method for learning the distributions of data and effectively bridged the gap between DL and Bayesian inference. This innovation marked the inception of a new class of generative models that has since been extensively used for a variety of applications including image generation and anomaly detection. VAEs are fundamentally different from traditional autoencoders. While traditional autoencoders aim to encode inputs into a deterministic latent representation, VAEs treat the latent space as probabilistic and model the latent variables as a distribution, typically Gaussian. This probabilistic approach allows VAEs to generate new data instances by sampling from the latent space, thus serving as generative models [40]. The encoding phase of a VAE involves mapping the input data $\mathbf{x}$ to a distribution over the latent space described by parameters (mean μ and variance σ^2), which are functions of the input data learned by the network. The latent variable $\mathbf{z}$ is then sampled from this distribution:

$$\mathbf{z} \sim \mathcal{N}(\mu(\mathbf{x}), \sigma(\mathbf{x})^2) \quad (12)$$

The decoder then reconstructs the input data from this sampled latent variable. This architecture is shown in Fig. 4. The entire model is trained by maximizing the likelihood of the observed data under the model, which is computationally intractable directly but facilitated by maximizing the evidence lower bound (ELBO) on the log-likelihood:

$$\mathcal{L}(\theta, \phi; \mathbf{x}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})}[\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) \quad (13)$$

where $q_{\phi}(\mathbf{z}|\mathbf{x})$ is the encoding distribution, $p_{\theta}(\mathbf{x}|\mathbf{z})$ is the decoding distribution, D_{KL} is the Kullback–Leibler divergence between the encoded distribution and the prior distribution $p(\mathbf{z})$, and θ, ϕ are parameters of the decoder and encoder, respectively. Since their introduction, VAEs have had a significant impact in the ML field, providing a framework for building complex generative models that can learn to produce new data instances that resemble the given training data. Their ability to generate new samples by varying the latent variables has been particularly beneficial in domains such as drug discovery, where generating new molecular structures is invaluable.


4.2.2 Conditional Autoencoders

Conditional autoencoders are an advanced variant of the standard autoencoder that incorporate additional contextual information into the autoencoding process [42]. By conditioning on auxiliary information such as class labels or other relevant data, these models are capable of generating or reconstructing outputs that are high in fidelity and tailored to specific conditions. The main idea behind conditional AEs is to modify both the encoder and decoder phases of the autoencoder by introducing additional conditional inputs. This approach allows the model to learn representations that are specific to given conditions, which can significantly enhance the relevance and utility of the learned encodings. The encoder function ϕ maps both the input $\mathbf{x}$ and the condition $\mathbf{y}$ to a latent representation $\mathbf{z}$, and the decoder function ψ then reconstructs the input from $\mathbf{z}$ under the same condition:

$$\mathbf{z} = \phi(\mathbf{x}, \mathbf{y}) \quad (14)$$

$$\hat{\mathbf{x}} = \psi(\mathbf{z}, \mathbf{y}) \quad (15)$$

Fig. 4 Basic variational autoencoder architecture [41]



This process ensures that the output $\hat{\mathbf{x}}$ is not only a reconstruction of $\mathbf{x}$ but is also customized according to $\mathbf{y}$ [43]. The condition $\mathbf{y}$ can be any relevant metadata, such as labels in supervised learning tasks, user information in recommendation systems, or any other auxiliary data that can help guide the reconstruction process. Conditional AEs have found extensive applications in domains where the condition can significantly influence the output. For example, in image processing, a conditional AE can generate images of a specific class when conditioned on class labels or produce images with particular attributes when conditioned on descriptors such as “smiling” or “wearing glasses” for facial images [44]. In the field of drug design, they can generate molecular structures conditioned on properties like solubility or toxicity [45]. The inclusion of conditions in the autoencoding process enhances the specificity of the generated outputs and improves the model’s ability to handle complex datasets where different subsets of data have distinct characteristics. This makes conditional AEs particularly useful in tasks requiring high degrees of customization and in scenarios where diverse outputs are needed from a single model.

4.2.3 Convolutional Autoencoders

Convolutional autoencoders integrate the robust features of convolutional neural networks (CNNs) into the autoencoder framework, making them especially effective for tasks involving image data. Through the utilization of convolutional layers, these models leverage the spatial hierarchy of features in images, which enhances their ability to learn efficient representations and perform high-quality reconstructions [46]. Unlike traditional autoencoders that use fully connected layers, convolutional AEs employ convolutional layers to encode input images into a compressed latent space [47, 48]. Convolutional layers apply a set of learnable filters to the input, which captures local feature patterns such as edges, textures, and other relevant image characteristics. After passing through several convolutional layers, typically

combined with pooling layers to reduce the spatial dimensions, the encoded data represents a compressed, feature-rich representation of the original input. The encoding process involves multiple convolutional layers where each layer applies a set of filters to extract higher-level features. These layers are often followed by activation functions such as the rectified linear unit (ReLU) to introduce non-linearity and pooling layers to reduce dimensionality and encode more abstract features into the latent space:

$$\mathbf{z} = \text{ConvPool}(\text{ReLU}(\text{Conv}(\mathbf{x}))) \quad (16)$$

The decoding process uses deconvolutional (i.e., transposed convolution) layers, which perform the reverse operation of convolutional layers [49]. Deconvolutional layers increase the spatial dimensions of the encoded representations, reconstructing the finer details and spatial structure of the original image as it progresses through each layer:

$$\hat{\mathbf{x}} = \text{Deconv}(\text{ReLU}(\text{DeconvPool}(\mathbf{z}))) \quad (17)$$

The architecture of a basic convolutional AE is shown in Fig. 5. Convolutional AEs are particularly valuable in image reconstruction, denoising, and anomaly detection in images [51–53]. They are also used in more sophisticated tasks like feature extraction for other complex ML models, image segmentation, and even in the generation of new images similar to those present in the training set. The main advantage of using convolutional AEs lies in their ability to efficiently handle image data. Lastly, they preserve the spatial structure of the input through convolutional operations, thereby reducing the loss of important visual information, leading to better performance in tasks requiring detailed image reconstruction or modification.

4.2.4 Sequence-to-Sequence Autoencoders

Sequence-to-sequence autoencoders are specialized architectures designed to handle sequential data such as text, audio,




Fig. 5 Basic architecture of convolutional autoencoder [50]

and time series [54]. They utilize the dynamic capabilities of recurrent neural networks (RNNs), or more often, Long Short-Term Memory (LSTM) or gated recurrent neural (GRU) networks to avoid the vanishing gradient problem, as shown in Fig. 6. These autoencoders effectively encode sequences into a fixed-size latent representation and then reconstruct the sequence from this representation, preserving temporal dependencies and sequence integrity [55]. The architecture of a Sequence-to-Sequence AE consists of two main components: an encoder and a decoder, both of which are typically implemented using RNNs. The encoder processes the input sequence, capturing the temporal information and compressing it into a latent state.

The decoder then takes this latent representation and generates an output sequence that aims to replicate the input sequence as closely as possible. During the encoding phase, the RNN processes the input sequence one element at a time, updating its internal state based on both the current input and its previous state [55]. This process continues until the entire sequence has been fed into the model, resulting in a final state that encapsulates the information from the entire sequence:

$$\mathbf{z} = \text{RNN}_{\text{encode}}(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n) \quad (18)$$

The decoding phase then begins, often initialized with a start symbol or the final state of the encoder. The decoder attempts to reconstruct the sequence, possibly predicting one step at a time and using its predictions as inputs for subsequent steps:

$$\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n = \text{RNN}_{\text{decode}}(\mathbf{z}) \quad (19)$$

Sequence-to-sequence AEs are incredibly versatile and find applications in numerous fields involving sequential data. They are particularly useful in natural language processing for tasks such as sentence compression, machine translation, and summarization [56, 57]. In the time series domain, they are employed for anomaly detection, forecasting, and signal processing [58]. The ability to model sequence data

effectively makes them ideal for these complex tasks. One of the key advantages of these type of AEs is their ability to handle variable-length input and output sequences, a feature particularly important in NLP tasks.

4.2.5 Contractive Autoencoders

Contractive autoencoders (CAEs) are a variant of traditional autoencoders designed to enhance the robustness and stability of the learned representations. They achieve this by incorporating a regularization term in the loss function, often referred to as the contractive term [59]. This term reduces the model's sensitivity to small changes in the input, thereby encouraging the autoencoder to learn more invariant and robust features. Specifically, the contractive term penalizes the Frobenius norm of the Jacobian matrix of the encoder's activations with respect to the input [60]. This penalty promotes a smooth encoding space where similar inputs produce similar encodings, effectively reducing the model's tendency to overfit to noise in the data. The loss function for a CAE includes the standard reconstruction loss plus the contractive term, enhancing the model's generalization capabilities:

$$L_{\text{CAE}} = L + \lambda \sum_i \|\nabla_{\mathbf{x}} h_i(\mathbf{x})\|^2 \quad (20)$$

where L denotes the reconstruction loss, such as mean squared error, λ is a regularization parameter that controls the weight of the contractive term, and $\nabla_{\mathbf{x}} h_i(\mathbf{x})$ represents the gradient of the i -th hidden unit with respect to the input vector $\mathbf{x}$. This mechanism is particularly valuable in tasks where stability and generalization over subtle input variations are critical. For example, in image recognition, variations in pixel intensity should not significantly alter the semantic content of the images. Thus, CAEs are extensively used in image processing and pattern recognition, where their ability to resist deformation and noise enhances performance. They are also beneficial in applications requiring feature




Fig. 6 Architecture of sequence-to-sequence autoencoder

extraction that remains invariant to minor local changes, supporting advanced tasks such as denoising, compression, and feature disentanglement.

4.2.6 Adversarial Autoencoders

AAEs are a sophisticated integration of autoencoder and generative adversarial network (GAN) technologies. By combining the autoencoding structure with the adversarial training concept from GANs, AAEs enforce a learned latent space that conforms to a specified prior distribution [11]. This approach improves the robustness of the encoding and enhances the generative capabilities of the autoencoder. In an AAE, the architecture is split into two main components: the autoencoder itself and an adversarial network, as shown in Fig. 7. The autoencoder comprises the typical encoder-decoder structure where the encoder compresses the input into a latent representation, and the decoder attempts to reconstruct the input from this representation. The adversarial network, often termed the discriminator, is trained to differentiate between the latent representations produced by the encoder and samples drawn from the desired prior distribution [61].

The training of an AAE occurs in two phases:

1. **Reconstruction Phase:** The autoencoder minimizes the reconstruction loss, typically measured as the mean squared error between the input and its reconstruction. This phase is similar to traditional autoencoder training and focuses on data compression and reconstruction capabilities.
2. **Regularization Phase:** Concurrently, the encoder also functions as a generator in a GAN setup, where it tries to fool the discriminator into believing that the encoded representations are valid samples drawn from the prior distribution. The discriminator is trained to distinguish between the true samples from the prior and the outputs of the encoder [61]. This adversarial process encourages the encoder to produce latent representations that closely mimic the prior.




Fig. 7 Architecture of adversarial autoencoder [61]

Furthermore, assuming L_{adv} is the adversarial loss, the reconstruction loss L_{recon} is represented mathematically as:

$$L = L_{recon} + \lambda L_{adv} \quad (21)$$

where λ is a weighting factor that balances the two losses. AAEs are particularly useful in applications requiring a controlled generation of data samples [11]. They are employed in tasks such as semi-supervised learning, where they leverage unlabeled data to improve learning from limited labeled data, and in anomaly detection, where the ability to model complex data distributions aids in identifying outliers. Additionally, AAEs are used in synthetic data generation where adherence to a specific distribution is crucial. The main advantage of AAEs over traditional autoencoders lies in their ability to explicitly control the distribution of the latent space, making them highly effective for generative tasks [62]. This controlled latent space also facilitates better interpretability and manipulation of latent variables, allowing for more precise modifications and generation of data instances. Table 1 summarises the various autoencoders.

5 Variational Autoencoders in Generative Modeling

Variational autoencoders have emerged as one of the foundational technologies in generative modeling, and they offer a powerful and flexible approach to learning and synthesizing complex data distributions. The basic VAE framework (discussed in Section 4.2.1) has been extended in various ways to enhance its generative capabilities and address some of its inherent limitations. These extensions are designed to improve the quality of the generation and increase the fidelity of the latent space and the model's ability to capture more complex data distributions. While VAEs have been a breakthrough in generative modeling, they are not without challenges. One of the most noted issues is their tendency to produce blurry samples, especially when compared to other generative models such as GANs [63]. This issue primarily arises from the model's inherent difficulty in capturing high-frequency details, a consequence of the Gaussian assumption in the latent space and the mean-squared error loss commonly used in reconstruction. Several advancements have been proposed to solve these problems, and they are discussed next.

5.1 Advancements with β -VAEs

The β -VAEs is a notable advancement aimed at overcoming the VAE's limitations. β -VAEs, an extension of the traditional VAE, introduce an adjustable hyperparameter β to

Table 1 Summary of autoencoder architectures

S/N	Architecture	Description	Pros	Cons
1	Sparse AE	Enhances feature selection by imposing sparsity in hidden layers	Promotes feature selectivity and robustness	Can lead to underfitting if too restrictive
2	Denoising AE	Trains to remove noise from corrupted inputs, enhancing feature robustness	Improves robustness to input noise	May require careful tuning of the noise level
3	Variational AE	Introduces probabilistic latent variables and generative capabilities	Enables generation of new data instances	Can suffer from blurry outputs
4	Conditional AE	Extends autoencoders with conditioning on additional information	Allows controlled output generation	Complexity increases with conditions
5	Convolutional AE	Uses convolutional layers, suitable for image data	Efficient for spatial data representation	Primarily limited to image-like data
6	Seq-to-Seq AE	Applies RNNs to encode and decode sequence data	Effective for temporal data like text and audio	Requires large datasets to perform well
7	Contractive AE	Adds a regularization term to make the model insensitive to small input changes	Encourages robust feature representation	Can be computationally intensive
8	Adversarial AE	Combines autoencoder with adversarial training to match a prior distribution	Provides a way to regularize the latent space	Training can be unstable

enhance the disentanglement of latent representations [64]. This modification aims to provide a more robust framework for balancing between the accuracy of the reconstruction and the abstraction of the latent space, which can be critical for certain applications like feature interpretation and manipulation. In β -VAEs, the hyperparameter β modifies the emphasis on the KL divergence term in the VAE's loss function. This term traditionally acts as a regularizer by forcing the learned distribution in the latent space to approximate a prior distribution, typically a standard Gaussian [65]. By adjusting β , practitioners can control how strictly this approximation is enforced:

$$\mathcal{L}_\beta = \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})] - \beta D_{KL}(q_\phi(\mathbf{z}|\mathbf{x}) \| p(\mathbf{z})) \quad (22)$$

A β value greater than 1 increases the weight of the KL divergence term, which encourages the encoder to learn more disentangled and less correlated features in the latent space, potentially at the cost of reconstructing the original inputs with lower fidelity [66]. Conversely, a β lower than 1 decreases this weight, potentially improving reconstruction accuracy but at the risk of reducing the disentanglement of the latent variables. β -VAEs have been particularly influential in tasks where the interpretability and manipulability of the latent space are paramount. For example, in facial recognition or style transfer applications, a more disentangled representation can allow for the independent manipulation of individual features such as adding a smile or changing the hairstyle in an image synthesis task [67, 68].

Studies have shown that increasing β leads to a latent space where changes along one dimension correspond to

changes in only one factor of variation, with little to no impact on other factors [69, 70]. This property is highly desirable in generative modeling, especially in creative and design applications where users need to adjust specific attributes without affecting others. While β -VAEs offer enhanced control over the trade-offs between reconstruction accuracy and feature disentanglement, they also require careful tuning of the β parameter. Overemphasis on disentanglement can lead to models that fail to capture enough detail to reconstruct inputs accurately, affecting their utility in applications where high fidelity is critical.

5.2 Enhancing Information Retention with InfoVAEs

InfoVAEs [71] tackle a significant challenge, i.e., the collapse of latent variables, where certain dimensions of the latent space become inactive or redundant. By maximizing the mutual information between the input and the latent variables, InfoVAEs ensure that the latent space retains meaningful information about the input data, enhancing the overall utility and expressiveness of the model. The core innovation in InfoVAEs is in their objective function, which aims to maximize the mutual information between the input data $\mathbf{x}$ and the latent variables $\mathbf{z}$, alongside minimizing the divergence from the prior distribution [71]. Mutual information is a measure of the amount of information obtained about one random variable through another random variable [72]. In the context of InfoVAEs, it quantifies how much information about the input $\mathbf{x}$ is captured by the latent variables $\mathbf{z}$:

$$\mathcal{L}_{\text{InfoVAE}} = \lambda \mathcal{I}(\mathbf{x}; \mathbf{z}) - D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) \quad (23)$$

where λ is a tuning parameter that balances the mutual information term against the KL divergence, D_{KL} , ensuring that the latent variables remain informative and relevant. InfoVAEs encourage the latent variables to retain as much information about the input data as possible by maximizing mutual information, preventing the model from overlooking critical features during the encoding process. InfoVAEs are useful in scenarios where each dimension of the latent space needs to convey meaningful and interpretable information about the data. This property is invaluable in fields such as bioinformatics, financial modeling, and other areas where understanding and interpreting the latent features can provide insights into complex patterns or relationships in the data.

While InfoVAEs provide significant advantages in terms of latent space utilization and feature expressiveness, they also require careful calibration of the λ parameter to balance between information retention and adherence to the prior distribution [73]. Misalignment in this balance can lead to either overfitting, where the model captures noise instead of signal, or underfitting, where the model fails to capture sufficient complexity from the data. Furthermore, VAEs have revolutionized generative modeling by providing a mechanism to learn complex data distributions and to generate new data samples that are statistically similar to the original dataset. This capability has opened new avenues in numerous fields ranging from image processing to financial modeling, where the generation of realistic samples can be critical for testing and simulation.

6 Recent Advancements and Methodologies

Autoencoders have evolved significantly in recent years, with advancements in neural network design and training methodologies driving substantial improvements in their performance and applicability. This section provides an overview of these recent innovations, highlighting how they have reshaped training practices and set new performance benchmarks in various domains.

6.1 Innovations in Autoencoder Design

6.1.1 Deep Autoencoders

The introduction of deeper architectures in autoencoder design marks a significant evolution, enabling these models to learn more complex and abstract data representations. Deep autoencoders, which utilize multiple layers of encoding and decoding, significantly enhance the capability to capture hierarchical features of the data. This progression

allows the models to handle tasks requiring detailed feature extraction. In deep autoencoder architectures, the input data is progressively compressed through successive layers of the encoder, each designed to capture different levels of data abstraction. At the deepest layer, the data representation is the most compressed, encapsulating the core characteristics of the input [40]. This representation is then progressively expanded through the decoder layers. Each decoder layer is typically designed to mirror a corresponding encoder layer, effectively reversing the encoding transformations to restore the original data features [74].

The benefits of incorporating depth into autoencoders are enormous. The additional layers allow the models to capture complex patterns in the data, such as distinguishing between various textures or objects within images, which are not possible with shallower architectures [21]. For instance, the first few layers might identify basic visual patterns like edges, while deeper layers might extract features related to the objects' interaction within the scene. Deep autoencoders have been beneficial in high-resolution image processing. They are used to enhance the quality of images in applications such as medical imaging, where preserving essential diagnostic details is crucial [20]. Furthermore, in the field of NLP, deep autoencoders enable the modeling of complex sentence structures, facilitating advanced applications like machine translation and sentiment analysis by understanding context and semantics over extended sequences.

Despite their advantages, training deep autoencoders is not without challenges. One major issue is the difficulty of training due to vanishing gradients, a phenomenon where gradients shrink as they backpropagate through the layers, making deeper layers harder to train effectively [75]. Techniques such as using unsupervised pre-training, batch normalization, and employing non-saturating activation functions like ReLUs have been developed to address these training challenges [76].

6.1.2 Adversarial Training in Autoencoders

Adversarial training is now a pivotal technique in the development of robust autoencoder architectures, enhancing their ability to generalize and resist adversarial attacks. This methodology involves training the autoencoder to withstand perturbations added to its input data, effectively simulating potential adversarial scenarios that the model might encounter in real-world applications [77]. The core idea behind adversarial training in autoencoders is to create a model that can recover clean data from inputs that have been intentionally corrupted with small but adversarially designed disturbances [61]. This process improves the model's robustness and also its ability to capture and reconstruct essential data features more accurately. For instance, adversarial training can be integrated into the

training process of VAEs or denoising autoencoders, where it assists in refining the model's loss, making them less sensitive to slight variations in input data that could lead to significant reconstruction errors.

Adversarial training typically involves alternating between training the autoencoder to minimize its usual loss function and adjusting it to maximize the effect of adversarial noise crafted by an adversarial network. This dual-training approach makes the autoencoder to be robust against potential manipulations or anomalies in data, thereby enhancing security and reliability, especially in sensitive applications like medical imaging and fraud detection.

6.1.3 Hybrid Models

Recent developments in neural network design have led to the integration of autoencoders with other sophisticated architectures, such as CNNs and RNNs [78–80]. This synthesis leverages the distinct advantages of each technique to significantly enhance the functionality and applicability of autoencoders across various domains. Hybrid models like Convolutional AE combine the autoencoder's data encoding and decoding capabilities with the spatial processing power of CNNs. This combination is useful for image data processing, where it exploits the spatial hierarchy of features—from basic edges and textures to more intricate patterns like entire objects—thus substantially improving the efficiency and precision of image reconstruction and analysis [80]. The ability of CNNs to handle spatial data complements the autoencoder's capacity for dimensionality reduction and feature encoding, making Convolutional AE invaluable for tasks such as medical imaging [81]. In these applications, they enhance image quality, facilitate anomaly detection, and reconstruct missing or corrupted parts of images, as illustrated by the work of [11], which discusses adversarial autoencoders that combine these principles.

Similarly, Sequence-to-Sequence Autoencoders incorporate RNNs to effectively manage sequential and temporal data. These models excel in domains requiring the processing of extended data sequences, such as text, speech, and time series. The research by [55] on sequence-to-sequence learning provides foundational insights into these applications, demonstrating the capability of RNNs to enhance language processing tasks when integrated within an autoencoder framework. The introduction of these hybrid models has not only expanded the scope of autoencoders but also created more versatile architectures capable of addressing a broader range of ML challenges. The convergence of different neural network methodologies within these models results in systems that are robust, accurate, and suitable for a diverse array of applications.

6.2 Emerging Methodologies in Training

6.2.1 Variational Techniques

The adoption of variational techniques in autoencoder training, like in VAEs, indicates a fundamental change in how these models are conceptualized and implemented. As described in Section 4.2.1., VAEs integrate the principles of Bayesian inference into autoencoder architectures, transforming the encoding process into a variational inference problem. This integration introduces a probabilistic approach to the training of autoencoders, where the encoder generates parameters of a proposed distribution over latent variables rather than deterministic outputs [65]. Instead of producing a latent space representation directly, the encoder outputs parameters of a probability distribution—typically Gaussian—defined by mean and variance. This stochastic approach allows the model to sample from this distribution to generate the latent representation, which is then used by the decoder to reconstruct the input [69]. Such a method not only helps in handling the inherent uncertainty in real-world data but also improves the robustness and generalizability of the model by effectively regularizing the latent space.

The flexibility and robustness provided by VAEs have broadened their applications in various fields such as image generation, financial modeling, and bioinformatics, where generating new, realistic samples and interpreting complex data distributions are crucial [82–84].

6.2.2 Self-Supervised Learning

Self-supervised learning is a rapidly growing area of research that applies principles of unsupervised learning to generate its training signals from the data itself. This approach has proven to be valuable in the training of autoencoders [85, 86]. It leverages the inherent structure of the data to create labels, which are then used as targets in training the autoencoder. In self-supervised learning, an autoencoder might be tasked with reconstructing the input data after some form of transformation, such as predicting the next frame in a video sequence or the missing part of an image. The autoencoder learns to predict these alterations based on the context provided by the input data, effectively training itself to understand and capture the underlying patterns and dependencies within the data. This technique assists in circumventing the need for large labeled datasets and enhances the model's ability to generalize from the data it has seen to new, unseen data. This capability is crucial for applications in fields such as computer vision, where models trained with self-supervised learning have been shown to achieve state-of-the-art performance on tasks like object detection and image segmentation [87, 88]. For example, a study by [89] demonstrated how predicting the relative position of

image patches could be used as a form of self-supervision for feature learning.

Similarly, in NLP, self-supervised learning techniques have enabled models to better understand the context and semantics of text, leading to improvements in language models and text generation systems. An example is the work by Devlin et al. [90] on BERT, which uses masked language modeling as a self-supervised task to enable deep language understanding. The implementation of self-supervised learning in autoencoder training continues to be a vibrant research area, with researchers developing increasingly sophisticated methods to extract more meaningful and complex patterns from unlabeled data.

6.2.3 Energy-Based Models

Energy-based models (EBMs) represent a compelling framework in ML, which conceptualizes learning as a process of energy minimization [91]. In recent years, there has been renewed interest in applying the principles of EBMs to autoencoders, leading to the development of energy-based autoencoders [92]. These models are designed to learn an energy function that assigns lower energy to observed or typical data points and higher energy to rare or anomalous data. The energy function typically measures the reconstruction error, with lower energy associated with smaller errors—indicating better reconstruction of the input data—and higher energy indicating significant deviations from the norm. By training the autoencoder to minimize the energy of clean data while maximizing the energy of anomalies or noise, the model effectively learns to filter out unwanted elements from the input data, which enhances the quality and reliability of the output. For instance, in anomaly detection, energy-based autoencoders can be trained on a dataset composed predominantly of normal occurrences. During operation, any input that significantly increases the reconstruction error, and consequently the energy, is flagged as anomalous [93]. This method has proven effective in various applications, from monitoring industrial systems [94, 95] to detecting fraudulent transactions [96], where anomalies may signify important or dangerous deviations from normal operations.

Similarly, in the field of image processing, energy-based autoencoders are used for denoising [97]. Through learning to assign high energy to noisy inputs and low energy to clean inputs, these models are able to reconstruct the original, noise-free image even when presented with significantly corrupted data. The ability to learn these energy distinctions makes them highly effective for applications requiring high levels of data integrity. Meanwhile, [98] presented the theoretical background of energy-based models in the context of autoencoders, and the study has been instrumental in developing and promoting the concept of energy-based learning in neural networks.

7 New Perspectives in Autoencoder Applications

Autoencoders have long been established in the domains of image and speech processing, but recent advances have pushed their applications into new areas and with novel data types. As datasets grow in complexity and diversity, autoencoders continue to evolve, meeting modern expectations with innovative solutions.

7.1 Expanding Horizons in Autoencoder Applications

Recent innovations have seen autoencoders being applied in less traditional yet critically important domains such as genomics, where they help in understanding complex genetic data, and in cybersecurity, where they are used for anomaly detection to identify potential threats. Moreover, autoencoders are now being increasingly utilized in the financial sector, assisting in risk assessment and fraud detection by learning to recognize patterns indicative of fraudulent behavior.

7.1.1 Genomics

In the field of genomics, autoencoders have been employed for tasks involving the processing and dimensional reduction of complex genetic data. The ability of autoencoders, especially VAEs, to distill high-dimensional data into meaningful, low-dimensional representations is instrumental in uncovering genetic markers associated with diseases [99]. This capability is crucial for advancing genetic research, speeding up processes, and enhancing the accuracy of identifying relevant genetic features. Autoencoders in genomics are vital in analyzing vast amounts of genetic information by encoding genes and their interactions into a compressed, interpretable format [100]. For example, VAEs have been applied to interpret complex biological data sets, allowing researchers to extract and visualize genetic features that are significant for understanding disease mechanisms [101].

A notable application involves using autoencoders for the identification of biomarkers in cancer genomics. Researchers employ these models to analyze thousands of genetic expressions to find those few that correlate strongly with specific cancer types [102, 103]. This method improves the targeted treatment approaches in personalized medicine, where treatments can be tailored based on the genetic profile of an individual's tumor. Furthermore, in drug discovery, autoencoders are utilized to predict how different genetic profiles will react to specific drug compounds [104]. This approach significantly aids in the development of new drugs by predicting efficacy and side effects, thus reducing the

costs and time involved in clinical trials. The use of autoencoders in genomics extends to more complex applications, such as predicting the progression of diseases and their possible mutations [105, 106]. By modeling genetic data over time, autoencoders can predict how a disease will evolve, which is vital for chronic diseases such as Alzheimer's or various forms of cancer.

7.1.2 Cybersecurity

Autoencoders have been used in cybersecurity, particularly in the area of anomaly detection. This capability is crucial for identifying unusual patterns and activities that may indicate security breaches or malicious attacks. Unlike traditional methods that often rely on predefined rules or signatures, autoencoders learn to model normal behavior from data traffic and user activities, enabling them to detect anomalies as deviations from this learned norm [107]. This adaptive approach allows autoencoders to stay effective even as cyber threats evolve with increasingly sophisticated tactics. In this field, autoencoders typically function by encoding network traffic or system logs into a compressed latent space and then reconstructing the output to match the input as closely as possible. The reconstruction error is a key indicator: a higher error suggests that the data deviates from the "normal" patterns the model has learned, indicating a potential anomaly or intrusion [108, 109]. This method is highly effective because it is not constrained by the limitations of predefined threat signatures, which may not capture new or evolving attack vectors.

Meanwhile, a notable study by Sakurada and Yairi [8] demonstrated the effectiveness of autoencoders in detecting anomalies in spacecraft data, highlighting the model's potential in high-stakes security scenarios. Autoencoders are also being used to enhance endpoint security. They can detect deviations that may indicate malware presence or data exfiltration attempts by monitoring applications and system processes. This application is valuable as endpoints are often the targets of initial compromise in larger cyber-attacks. Furthermore, in the context of user behavior analytics, autoencoders have been used in profiling normal user behavior and identifying activities that stray from these profiles, such as unusual login times or uncharacteristic access to sensitive resources [110]. These capabilities make autoencoders an integral part of sophisticated cybersecurity systems aiming to preemptively identify and mitigate potential security threats.

7.1.3 Financial Services

Autoencoders have become an integral tool in financial services, particularly for detecting anomalies in transaction data that could indicate fraudulent activities. They can

effectively identify and flag transactions that deviate significantly from established norms by learning what constitutes "normal" transaction patterns [111]. This capability is crucial in areas such as real-time payment processing, credit card fraud detection, and credit scoring, where speed, accuracy, and the ability to adapt to new types of fraudulent tactics are paramount. In the context of fraud detection, AEs are trained on vast amounts of transaction data, learning to compress this data into a lower-dimensional latent space and then reconstructing it. The reconstruction error serves as a critical metric, with transactions that result in higher-than-typical errors being flagged for further investigation [112]. This method allows financial institutions to detect a wide range of anomalous behaviors automatically, from simple outliers to complex fraud models that might not be evident through traditional detection methods.

One of the key advantages of using autoencoders in financial fraud detection is their ability to adapt over time. As new transaction data is continuously fed into the system, the autoencoder can refine its understanding of what normal transaction patterns look like, thus staying effective even as fraudsters evolve their methods [113]. Moreover, autoencoders have been applied in credit scoring models where they assess the creditworthiness of applicants by analyzing their transaction histories and financial behaviors [83, 114]. By identifying patterns that correlate with risky behaviors, financial institutions can better manage their risk profiles and make more informed lending decisions. A notable application of autoencoders in financial services is highlighted in the research by Poutré et al. [115], which explores the use of these models for anomaly detection in financial market data. This study shows the potential of autoencoders not only in detecting fraudulent transactions but also in identifying market manipulation and other irregular financial activities.

7.2 Adapting to Modern Datasets and Expectations

As datasets grow in size and complexity, the inherent flexibility and scalability of autoencoders make them indispensable. Their ability to compress data into a condensed latent space enables detailed yet manageable analyses of large datasets. This section explores their ability to adapt to modern datasets.

7.2.1 Integration with ML Pipelines

Autoencoders are increasingly seen as vital components within broader ML pipelines, particularly for their ability to serve as effective pre-training mechanisms. By initially training on large volumes of unlabeled data, autoencoders learn to capture the essential structures and patterns within that data, effectively encoding this information into a lower-dimensional space. This process not only reduces the

dimensionality of the data but also filters out noise and less relevant information, thereby enhancing the quality of features available for subsequent learning stages. When these pre-trained models are then used as the starting point for supervised learning tasks, they significantly improve the efficiency and accuracy of the resulting models. For instance, in fields like image recognition and speech analysis, autoencoders pre-trained on vast datasets can significantly reduce the amount of labeled data required to achieve high levels of accuracy [116]. Furthermore, the features learned by autoencoders can often reveal new insights about the data, leading to more robust and generalizable models. The ability of autoencoders to pre-process and transform data makes them a natural fit for integration into ML workflows, particularly in an era where data complexity and volume pose significant challenges. Through simplifying data and extracting relevant features without the need for labels, autoencoders allow subsequent models to focus on refining predictions rather than understanding the raw data from scratch.

7.2.2 Handling Complex and Unstructured Data

Autoencoders have demonstrated remarkable adaptability to various types of data. This adaptability is important when dealing with unstructured data, which constitutes a significant portion of the data generated today—from social media content to multimedia and sensor data. Unstructured data, such as images, video, text, and audio, does not fit neatly into traditional database formats [117]. Autoencoders simplify the complexity of processing such data by learning efficient representations that are often much smaller in size yet retain most of the critical information. For example, in image and video processing, autoencoders reduce data dimensionality, enabling easier storage, processing, and analysis. They effectively capture the underlying patterns in visual content, which can be used for tasks like image segmentation, object recognition, and more sophisticated video analytics [118, 119].

Furthermore, in text data, autoencoders transform lengthy documents into dense vectors that encapsulate the semantic properties of the words and sentences [120]. This capability is very useful for natural language processing applications such as topic modeling, sentiment analysis, and automated summarization. By extracting and learning from patterns in customer feedback or social media discussions, businesses can gain insights into public sentiment and emerging trends without the need for manually tagging or categorizing the content. Additionally, autoencoders are employed to analyze social media content, enabling businesses to understand and predict customer behavior and preferences [121]. By learning from vast amounts of user-generated content, autoencoders can assist marketers in identifying patterns that indicate customer interests, engagement levels, and potential areas

for product development or marketing strategies [122]. The flexibility of autoencoders also extends to sensor data and Internet of Things (IoT) applications, where they are used to monitor and analyze data from various sensors and devices [123]. In these settings, autoencoders help detect anomalies, predict equipment failures, or optimize energy usage in smart grids by learning normal operational patterns and identifying deviations.

7.3 Autoencoders as Generative Models

Autoencoders, traditionally used for dimensionality reduction and feature extraction, have increasingly been explored for their capabilities as generative models. This shift is driven by their ability to learn detailed representations of data, which can then be manipulated to generate new, realistic samples that mimic the original data distribution [124]. Autoencoders, particularly through innovations such as VAEs and AAEs, have extended their utility beyond mere data compression to become powerful generative models. These models have adapted the foundational autoencoder architecture to not only reconstruct inputs but also generate new data instances that adhere to the learned distribution of the input data. VAEs present an approach that allows for the generation of new data points by sampling different combinations of latent variables, thus producing diverse and novel outputs [125]. The flexibility in generation comes from the stochastic nature of the latent space representation, which encapsulates the data's underlying distribution. This aspect was foundational in the work by Kingma and Welling [10], who demonstrated that such a setup not only helps in effective data reconstruction but also in generating new data that maintain statistical properties similar to those of the training set.

Meanwhile, AAEs leverage the adversarial training concept introduced by Goodfellow et al. [12] in the development of GANs. In AAEs, the autoencoder architecture is paired with a discriminator that tries to distinguish between encoded representations and samples from a prior distribution. This adversarial process compels the autoencoder to produce latent representations that are indistinguishable from the prior, leading to a more regularized and diverse generative process. This model has shown particular effectiveness in applications where maintaining the diversity of output is crucial, such as in generating realistic images or complex sequences of data. Both VAEs and AAEs have significantly impacted fields requiring robust generative models, including drug discovery, where they generate novel molecular structures [126], and in creative industries, where they contribute to the arts by creating music, art, and literary works [127]. Their ability to produce new content that follows the learned patterns and complexities of substantial datasets has made them invaluable in diverse applications.

Furthermore, autoencoders have been used in image generation to produce new images that share characteristics with a training set, such as generating new facial images from a dataset of faces [128]. This application is not only useful for artistic purposes but also for augmenting datasets in ML, where more training data can lead to improved model performance. In text generation, autoencoders learn to encode textual information in a latent space, and manipulation of this space allows for the generation of new text sequences that are coherent and contextually relevant, showing promise in applications such as chatbots and automated story generation.

8 Future Research Directions

Autoencoders have evolved from simple dimensionality reduction tools to sophisticated models capable of handling a wide range of data processing tasks. Their application in domains such as genomics, cybersecurity, and financial services has demonstrated their versatility and efficacy, as discussed in Section 7. Particularly, VAEs and AAEs have pushed the boundaries of what autoencoders can achieve, transitioning them into powerful generative models that can produce realistic and diverse outputs. Despite their advancements, autoencoders face several challenges that limit their efficiency and applicability, and future research will focus on such challenges.

One of the significant challenges in the field of generative autoencoders, particularly VAEs and AAEs, is the tendency to produce overly smooth or unrealistic outputs [129, 130]. This issue often stems from limitations in how these models construct their latent space representations and from the constraints of their decoding mechanisms. To enhance the fidelity of generated outputs, further research is needed to refine these aspects of autoencoder design. Enhancing latent space representations could involve integrating more complex prior distributions within VAEs. Traditional VAEs typically use simple Gaussian distributions to model the latent space, which may be insufficient for capturing the multi-modal and complex nature of many real-world data distributions. By exploring alternative distributions, such as mixture models or heavy-tailed distributions, researchers can potentially enable VAEs to model data with higher accuracy and greater detail, thereby producing more realistic and varied outputs.

Another promising avenue is the use of advanced hybrid models that incorporate the strengths of different neural network architectures alongside autoencoders. For instance, combining autoencoders with CNNs for image data or with RNNs for sequence data can leverage the specialized capabilities of these architectures to improve the quality of the reconstruction [57, 131]. CNNs can add robust spatial

hierarchies to the encoding-decoding process, essential for tasks like image and video generation, while RNNs can contribute their temporal dynamics modeling capabilities, which are crucial for generating realistic sequences of text or speech. Furthermore, improving the decoding mechanisms of autoencoders also holds potential. Developing decoders that can better reconstruct the complexity and variability of input data is crucial. This might involve innovations in network design, such as using more complex activation functions, deeper network architectures, or novel network training techniques like attention mechanisms or capsule networks, which have shown promise in enhancing the representational capabilities of neural networks. Ongoing research into the integration of external memory systems or the incorporation of reinforcement learning strategies to guide the generative process could also enhance output diversity and realism. These approaches can provide additional context or guidance to the generative model, allowing for more controlled and accurate generation processes.

Furthermore, with their inherent ability to learn useful representations from unlabeled data, autoencoders hold significant potential in unsupervised and semi-supervised learning contexts. By exploring novel ways to leverage autoencoders in these contexts, future studies could unlock new avenues for applying ML to broader sets of problems, especially in domains where acquiring labeled data poses a significant challenge. In unsupervised learning, autoencoders can be used to discover the underlying structure and distribution of data without any labeled examples. This capability is crucial for tasks like clustering, dimensionality reduction, and anomaly detection. Future research could focus on developing more sophisticated autoencoder architectures that are better at capturing complex and high-level abstractions in data. For instance, introducing modular or hierarchical autoencoders could allow for learning at multiple scales and granularities, which could significantly enhance the quality and usefulness of the learned representations.

Another area for expanding the applications of autoencoders in unsupervised and semi-supervised contexts is through transfer learning. Autoencoders trained on one task or dataset can potentially be adapted to other tasks or datasets, leveraging the generality of the learned features. This approach can be effective in domains with similar data structures but where labeled data is only available for specific tasks. Research into how autoencoder-based models can be efficiently transferred across different domains without extensive retraining could significantly broaden their applicability and effectiveness.

Finally, the development of continuous learning systems, where autoencoders continually adapt and update their learned representations based on new data, represents an exciting frontier. Such systems could be critical for dynamic environments where data distributions change over time,

such as in financial markets or evolving social networks. Research into continuous or lifelong learning models would address one of the fundamental challenges in current ML paradigms, which is the ability to adapt to new information without forgetting previously learned knowledge. Expanding the applications of autoencoders in unsupervised and semi-supervised learning could revolutionize how we approach ML in data-rich yet label-scarce environments. By enhancing feature learning, improving learning efficiencies, and developing adaptable systems, autoencoders could significantly influence a wide range of industries, from healthcare to finance, by providing more accurate, robust, and adaptable ML solutions.

9 Conclusion

Autoencoders have proven to be versatile and powerful in diverse applications, advancing beyond their initial applications in dimensionality reduction to address a wide range of complex and data-intensive challenges across various domains. This study presented a comprehensive review of autoencoder neural networks and their evolution from the basic architectures to the recent state-of-the-art variational and adversarial autoencoders. The study has highlighted their evolving roles, particularly emphasizing their capabilities as generative models, their utility in unsupervised and semi-supervised learning environments, and their expanding influence in cross-disciplinary applications.

Funding Open access funding provided by University of Johannesburg.

Open Access This article is licensed under a Creative Commons Attribution 4.0 International License, which permits use, sharing, adaptation, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if changes were made. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit <http://creativecommons.org/licenses/by/4.0/>.

References

- Li Y, Yang M, Zhang Z (2018) A survey of multi-view representation learning. *IEEE Trans Knowl Data Eng* 31(10):1863–1883
- Zhang C, Yang Z, He X, Deng L (2020) Multimodal intelligence: representation learning, information fusion, and applications. *IEEE J Select Topics Signal Process* 14(3):478–493
- Schölkopf B, Locatello F, Bauer S, Ke NR, Kalchbrenner N, Goyal A, Bengio Y (2021) Toward causal representation learning. *Proc IEEE* 109(5):612–634
- Bengio Y, Courville A, Vincent P (2013) Representation learning: a review and new perspectives. *IEEE Trans Pattern Anal Mach Intell* 35(8):1798–1828
- Mienye ID, Sun Y (2021) Improved heart disease prediction using particle swarm optimization based stacked sparse autoencoder. *Electronics* 10(19):2347
- Charte D, Charte F, Jesus MJ, Herrera F (2020) An analysis on the use of autoencoders for representation learning: fundamentals, learning task case studies, explainability and challenges. *Neurocomputing* 404:93–107
- Vincent P, Larochelle H, Bengio Y, Manzagol P-A (2008) Extracting and composing robust features with denoising autoencoders. In: *Proceedings of the 25th international conference on machine learning*, pp. 1096–1103
- Sakurada M, Yairi T (2014) Anomaly detection using autoencoders with nonlinear dimensionality reduction. In: *Proceedings of the 2014 2nd workshop on machine learning for sensory data analysis*, pp. 4–11
- Hinton GE, Salakhutdinov RR (2006) Reducing the dimensionality of data with neural networks. *Science* 313(5786):504–507
- Kingma DP, Welling M (2013) Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*
- Makhzani A, Shlens J, Jaitly N, Goodfellow I, Frey B (2015) Adversarial autoencoders. *arXiv preprint arXiv:1511.05644*
- Goodfellow I, Pouget-Abadie J, Mirza M, Xu B, Warde-Farley D, Ozair S, Courville A, Bengio Y (2014) Generative adversarial nets. *Advances in Neural Information Processing Systems* 27
- Qi G-J, Luo J (2020) Small data challenges in big data era: a survey of recent progress on unsupervised and semi-supervised methods. *IEEE Trans Pattern Anal Mach Intell* 44(4):2168–2187
- Baldi P (2012) Autoencoders, unsupervised learning, and deep architectures. In: *Proceedings of ICML workshop on unsupervised and transfer learning*, pp. 37–49. *JMLR workshop and conference proceedings*
- Tschannen M, Bachem O, Lucic M (2018) Recent advances in autoencoder-based representation learning. *arXiv preprint arXiv:1812.05069*
- Pinaya WHL, Vieira S, Garcia-Dias R, Mechelli A (2020) Autoencoders. *Machine Learning*. Elsevier, pp 193–208
- Ng A, et al. (2011) Sparse autoencoder. *CS294A Lecture notes* 72, 1–19
- Nayak R, Pati UC, Das SK (2021) A comprehensive review on deep learning-based methods for video anomaly detection. *Image Vis Comput* 106:104078
- Qian J, Song Z, Yao Y, Zhu Z, Zhang X (2022) A review on autoencoder based representation learning for fault detection and diagnosis in industrial processes. *Chemom Intell Lab Syst* 231:104711
- Mienye ID, Sun Y, Wang Z (2020) Improved predictive sparse decomposition method with densenet for prediction of lung cancer. *Int J Comput* 1:533–541
- Shang Z, Sun L, Xia Y, Zhang W (2021) Vibration-based damage detection for bridges by deep convolutional denoising autoencoder. *Struct Health Monit* 20(4):1880–1903
- Bank D, Koenigstein N, Giryas R (2023) Autoencoders. In: *Machine learning for data science handbook: data mining and knowledge discovery handbook*, pp 353–374
- Shi H, Yang N, Tang H, Yang X (2022) Asgd: stochastic gradient descent with adaptive batch size for every parameter. *Mathematics* 10(6):863
- Mohtashami A, Jaggi M, Stich SU (2023) Special properties of gradient descent with large learning rates. *International conference on machine learning*. PMLR, pp 25082–25104

25. Tan HH, Lim KH (2019) Vanishing gradient mitigation with deep learning neural network optimization. 2019 7th international conference on smart computing & communications. IEEE, pp 1–4
26. Jaiswal G, Rani R, Mangotra H, Sharma A (2023) Integration of hyperspectral imaging and autoencoders: benefits, applications, hyperparameter tuning and challenges. *Comput Sci Rev* 50:100584
27. Hinton G.E, Zemel R (1993) Autoencoders, minimum description length and Helmholtz free energy. *Adv Neural Inf Process Syst* 6
28. Van Der Maaten L, Postma E, Herik J et al (2009) Dimensionality reduction: a comparative. *J Mach Learn Res* 10:66–71
29. Friedrich T (2002) Nonlinear dimensionality reduction with locally linear embedding and isomap. Masters Dissertation, University of Sheffield
30. Kramer MA (1991) Nonlinear principal component analysis using autoassociative neural networks. *AICHE J* 37(2):233–243
31. Fan Z, Bi D, He L, Shiping M, Gao S, Li C (2017) Low-level structure feature extraction for image processing via stacked sparse denoising autoencoder. *Neurocomputing* 243:12–20
32. Chandar AP S, Lauly S, Larochelle H, Khapra M, Ravindran B, Raykar V.C, Saha A (2014) An autoencoder approach to learning bilingual word representations. *Adv Neural Inf Process Syst* 27
33. Olshausen BA, Field DJ (1996) Natural image statistics and efficient coding. *Netw Comput Neural Syst* 7(2):333
34. Steck H (2019) Embarrassingly shallow autoencoders for sparse data. In: The World Wide Web conference, pp 3251–3257
35. Hoeffer T, Alistarh D, Ben-Nun T, Dryden N, Peste A (2021) Sparsity in deep learning: pruning and growth for efficient inference and training in neural networks. *J Mach Learn Res* 22(241):1–124
36. Liang J, Liu R (2015) Stacked denoising autoencoder and dropout together to prevent overfitting in deep neural network. In: 2015 8th international congress on image and signal processing (CISP), pp 697–701. <https://doi.org/10.1109/CISP.2015.7407967>
37. El-Shafai W, El-Nabi S.A, Ali AM, El-Rabaie E-SM, Abd El-Samie FE (2023) Traditional and deep-learning-based denoising methods for medical images. *Multimedia Tools App*, pp 1–28
38. Lu X, Tsao Y, Matsuda S, Hori C (2013) Speech enhancement based on deep denoising autoencoder. *Interspeech* 2013:436–440
39. Abouzid H, Chakkor O, Reyes OG, Ventura S (2019) Signal speech reconstruction and noise removal using convolutional denoising autoencoders with neural deep learning. *Analog Integr Circ Sig Process* 100:501–512
40. Kingma DP, Welling M et al (2019) An introduction to variational autoencoders. *Foundations Trends(R) Mach Learn* 12(4):307–392
41. Joshi M, Pant DR, Karn RR, Heikkonen J, Kanth R (2022) Meta-learning, fast adaptation, and latent representation for head pose estimation. In: 2022 31st conference of open innovations association (FRUCT), ppD 71–78. <https://doi.org/10.23919/FRUCT54823.2022.9770932>
42. Lopez-Martin M, Carro B, Sanchez-Esguevillas A, Lloret J (2017) Conditional variational autoencoder for prediction and feature recovery applied to intrusion detection in iot. *Sensors* 17(9):1967
43. Choi Y, El-Khamy M, Lee J (2019) Variable rate deep image compression with a conditional autoencoder. In: Proceedings of the IEEE/CVF international conference on computer vision, pp 3146–3154
44. Li H-Y, Dong W-M, Hu B-G (2018) Facial image attributes transformation via conditional recycle generative adversarial networks. *J Comput Sci Technol* 33:511–521
45. Sousa T, Correia J, Pereira V, Rocha M (2021) Generative deep learning for targeted compound design. *J Chem Inf Model* 61(11):5343–5361
46. Mou L, Ghamisi P, Zhu XX (2017) Unsupervised spectral-spatial feature learning via deep residual conv-deconv network for hyperspectral image classification. *IEEE Trans Geosci Remote Sens* 56(1):391–406
47. La Grassa R, Re C, Cremonese G, Gallo I (2022) Hyperspectral data compression using fully convolutional autoencoder. *Remote Sens* 14(10):2472
48. Cheng Z, Sun H, Takeuchi M, Katto J (2019) Energy compaction-based image compression using convolutional autoencoder. *IEEE Trans Multimedia* 22(4):860–873
49. Ghosh SK, Biswas B, Ghosh A (2019) Sdca: a novel stack deep convolutional autoencoder-an application on retinal image denoising. *IET Image Proc* 13(14):2778–2789
50. Kanagamani T, Chakravarthy VS, Ravindran B, Menon RN (2023) A deep network-based model of hippocampal memory functions under normal and Alzheimer's disease conditions. *Front Neural Circ* 17:1092933
51. Wang S, Wang X, Zhang L, Zhong Y (2021) Auto-ad: autonomous hyperspectral anomaly detection network based on fully convolutional autoencoder. *IEEE Trans Geosci Remote Sens* 60:1–14
52. Maggipinto M, Beghi A, Susto GA (2022) A deep convolutional autoencoder-based approach for anomaly detection with industrial, non-images, 2-dimensional data: a semiconductor manufacturing case study. *IEEE Trans Autom Sci Eng* 19(3):1477–1490
53. Lin D, Li Y, Xie S, Nwe TL, Dong S (2023) DDR-ID: dual deep reconstruction networks based image decomposition for anomaly detection. *J Ambient Intell Humaniz Comput* 14(3):2125–2139
54. Ghimire S, Deo RC, Wang H, Al-Musaylh MS, Casillas-Pérez D, Salcedo-Sanz S (2022) Stacked LSTM sequence-to-sequence autoencoder with feature selection for daily solar radiation prediction: a review and new modeling results. *Energies* 15(3):1061
55. Sutskever I, Vinyals O, Le QV (2014) Sequence to sequence learning with neural networks. *Adv Neural Inf Process Syst* 27
56. Baziotis C, Androutsopoulos I, Konstas I, Potamianos A (2019) Seq 3: differentiable sequence-to-sequence-to-sequence autoencoder for unsupervised abstractive sentence compression. *arXiv preprint arXiv:1904.03651*
57. Jang M, Seo S, Kang P (2019) Recurrent neural network-based semantic variational autoencoder for sequence-to-sequence learning. *Inf Sci* 490:59–73
58. Nam KH, Song YJ, Yun ID (2021) SSS-AE: anomaly detection using self-attention based sequence-to-sequence auto-encoder in SMD assembly machine sound. *IEEE Access* 9:131191–131202
59. Diallo B, Hu J, Li T, Khan GA, Liang X, Zhao Y (2021) Deep embedding clustering based on contractive autoencoder. *Neurocomputing* 433:96–107
60. Rifai S, Bengio Y, Dauphin Y, Vincent P (2012) A generative process for sampling contractive auto-encoders. *arXiv preprint arXiv:1206.6434*
61. Pinho E, Costa C (2018) Feature learning with adversarial networks for concept detection in medical images. In: *UA.PTBioinformatics at ImageCLEF 2018*
62. Pidhorskyi S, Adjero DA, Doretto G (2020) Adversarial latent autoencoders. In: *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp 14104–14113
63. Kos J, Fischer I, Song D (2018) Adversarial examples for generative models. 2018 IEEE security and privacy workshops (SPW). IEEE, pp 36–42
64. Sikka H, Zhong W, Yin J, Pehlevant C (2019) A closer look at disentangling in β -vae. 2019 53rd Asilomar conference on signals, systems, and computers. IEEE, pp 888–895
65. Asperti A, Trentin M (2020) Balancing reconstruction error and Kullback–Leibler divergence in variational autoencoders. *IEEE Access* 8:199440–199448

66. Eddahmani I, Pham C-H, Napoléon T, Badoc I, Fouefack J-R, El-Bouz M (2023) Unsupervised learning of disentangled representation via auto-encoding: a survey. *Sensors* 23(4):2362
67. Zhang J, Huang Y, Li Y, Zhao W, Zhang L (2019) Multi-attribute transfer via disentangled representation. *Proc AAAI Conf Artif Intell* 33:9195–9202
68. Nitzan Y, Bermano A, Li Y, Cohen-Or D (2020) Face identity disentanglement via latent space mapping. *arXiv preprint arXiv:2005.07728*
69. Mathieu E, Rainforth T, Siddharth N, Teh YW (2019) Disentangling disentanglement in variational autoencoders. *International conference on machine learning*. PMLR, pp 4402–4412
70. Higgins I, Matthey L, Pal A, Burgess CP, Glorot X, Botvinick MM, Mohamed S, Lerchner A (2017) Beta-vae: Learning basic visual concepts with a constrained variational framework. *Int Conf Learn Representations (Poster)* 3
71. Zhao S, Song J, Ermon S (2018) InfoVAE: information maximizing variational autoencoders
72. Kraskov A, Stögbauer H, Grassberger P (2004) Estimating mutual information. *Phys Rev E* 69(6):066138
73. Rybkin O, Daniilidis K, Levine S (2021) Simple and effective VAE training with calibrated decoders. *International conference on machine learning*. PMLR, pp 9179–9189
74. Xie Y, Thurey N (2023) Reviving autoencoder pretraining. *Neural Comput Appl* 35(6):4587–4619
75. Liu D, Wang Y, Luo C, Ma J (2023) An improved autoencoder for recommendation to alleviate the vanishing gradient problem. *Knowl-Based Syst* 263:110254
76. Bengio Y et al (2009) Learning deep architectures for AI. *Foundations and trends® in Machine Learning* 2(1):1–127
77. Yuan X, He P, Zhu Q, Li X (2019) Adversarial examples: attacks and defenses for deep learning. *IEEE Trans Neural Netw Learn Syst* 30(9):2805–2824
78. Jana D, Patil J, Herkal S, Nagarajaiah S, Duenas-Orsorio L (2022) CNN and convolutional autoencoder (CAE) based real-time sensor fault detection, localization, and correction. *Mech Syst Signal Process* 169:108723
79. Ullah A, Muhammad K, Haq IU, Baik SW (2019) Action recognition using optimized deep autoencoder and CNN for surveillance data streams of non-stationary environments. *Futur Gener Comput Syst* 96:386–397
80. Dastider AG, Sadik F, Fattah SA (2021) An integrated autoencoder-based hybrid CNN-LSTM model for Covid-19 severity prediction from lung ultrasound. *Comput Biol Med* 132:104296
81. Badža MM, Barjaktarović MČ (2021) Segmentation of brain tumors from MRI images using convolutional autoencoder. *Appl Sci* 11(9):4317
82. Singh A, Ogunfunmi T (2021) An overview of variational autoencoders for source separation, finance, and bio-signal applications. *Entropy* 24(1):55
83. Mancisidor RA, Kampffmeyer M, Aas K, Jenssen R (2021) Learning latent representations of bank customers with the variational autoencoder. *Expert Syst Appl* 164:114020
84. Wang Z, Ren Y, Zhang X, Wang Y (2024) Generating long financial report using conditional variational autoencoders with knowledge distillation. *IEEE Trans Artif Intell*
85. Nousi P, Tefas A (2020) Self-supervised autoencoders for clustering and classification. *Evol Syst* 11(3):453–466
86. Gatopoulos I, Tomczak JM (2021) Self-supervised variational auto-encoders. *Entropy* 23(6):747
87. Ziegler A, Asano YM (2022) Self-supervised learning of object parts for semantic segmentation. In: *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp 14502–14511
88. Pan X, Tang F, Dong W, Gu Y, Song Z, Meng Y, Xu P, Deussen O, Xu C (2020) Self-supervised feature augmentation for large image object detection. *IEEE Trans Image Process* 29:6745–6758
89. Doersch C, Gupta A, Efros AA (2015) Unsupervised visual representation learning by context prediction. In: *Proceedings of the IEEE international conference on computer vision*, pp 1422–1430
90. Devlin J, Chang M-W, Lee K, Toutanova K (2018) Bert: pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*
91. Lv P, Yu Y, Fan Y, Tang X, Tong X (2020) Layer-constrained variational autoencoding kernel density estimation model for anomaly detection. *Knowl-Based Syst* 196:105753
92. Han T, Nijkamp E, Zhou L, Pang B, Zhu S-C, Wu YN (2020) Joint training of variational auto-encoder and latent energy-based model. In: *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 7978–7987
93. Basora L, Olive X, Dubot T (2019) Recent advances in anomaly detection methods applied to aviation. *Aerospace* 6(11):117
94. Kaymakci C, Wenninger S, Sauer A (2021) Energy anomaly detection in industrial applications with long short-term memory-based autoencoders. *Procedia CIRP* 104:182–187
95. Deng Z, Li Y, Zhu H, Huang K, Tang Z, Wang Z (2020) Sparse stacked autoencoder network for complex system monitoring with industrial applications. *Chaos Solitons Fractals* 137:109838
96. Misra S, Thakur S, Ghosh M, Saha SK (2020) An autoencoder based model for detecting fraudulent credit card transaction. *Procedia Comput Sci* 167:254–262
97. Li Z, Chen Y, Sommer FT (2023) Learning energy-based models in high-dimensional spaces with multiscale denoising-score matching. *Entropy* 25(10):1367
98. LeCun Y, Chopra S, Hadsell R, Ranzato M, Huang F (2006) A tutorial on energy-based learning. *Predicting Structured Data* 1
99. Wu Q, Bouez A, Bozkurt A, Masoomi A, Wang A, DeMeo DL, Weiss ST, Qiu W (2018) Deep learning methods for predicting disease status using genomic data. *J Biometrics Biostat* 9:5
100. Lopez R, Regier J, Cole MB, Jordan MI, Yosef N (2018) Deep generative modeling for single-cell transcriptomics. *Nat Methods* 15(12):1053–1058
101. Simidjievski N, Bodnar C, Tariq I, Scherer P, Andres Terre H, Shams Z, Jamnik M, Liò P (2019) Variational autoencoders for cancer data integration: design principles and computational practice. *Front Genet* 10:1205
102. Xie R, Wen J, Quitadamo A, Cheng J, Shi X (2017) A deep auto-encoder model for gene expression prediction. *BMC Genomics* 18:39–49
103. Gokhale M, Mohanty SK, Ojha A (2022) A stacked autoencoder based gene selection and cancer classification framework. *Biomed Signal Process Control* 78:103999
104. Li X, Xu Y, Lai L, Pei J (2018) Prediction of human cytochrome p450 inhibition using a multitask deep autoencoder neural network. *Mol Pharm* 15(10):4336–4345
105. Kmetzsch V, Becker E, Saracino D, Rinaldi D, Camuzat A, Le Ber I, Colliot O, et al (2022) Disease progression score estimation from multimodal imaging and microRNA data using supervised variational autoencoders. *IEEE J Biomed Health Inf* 26(12):6024–6035
106. Way GP, Greene CS (2018) Extracting a biologically relevant latent space from cancer transcriptomes with variational autoencoders. *Proceedings of the Pacific symposium on biocomputing*. World Scientific, pp 80–91
107. Said Elsayed M, Le-Khac N-A, Dev S, Jurcut AD (2020) Network anomaly detection using lstm based autoencoder. In: *Proceedings of the 16th ACM symposium on QoS and security for wireless and mobile networks*, pp 37–45

108. Torabi H, Mirtaheri SL, Greco S (2023) Practical autoencoder based anomaly detection by using vector reconstruction error. *Cybersecurity* 6(1):1
109. Xu W, Jang-Jaccard J, Singh A, Wei Y, Sabrina F (2021) Improving performance of autoencoder-based network anomaly detection on NSL-KDD dataset. *IEEE Access* 9:140136–140146
110. Deng S, Cai Q, Zhang Z, Wu X (2021) User behavior analysis based on stacked autoencoder and clustering in complex power grid environment. *IEEE Trans Intell Transp Syst* 23(12):25521–25535
111. Rubio J, Barucca P, Gage G, Arroyo J, Morales-Resendiz R (2020) Classifying payment patterns with artificial neural networks: an autoencoder approach. *Latin Am J Central Bank* 1(1–4):100013
112. Koo K, Park M, Yoon B (2024) A suspicious financial transaction detection model using autoencoder and risk-based approach. *IEEE Access*
113. Hilal W, Gadsden SA, Yawney J (2022) Financial fraud: a review of anomaly detection techniques and recent advances. *Expert Syst Appl* 193:116429
114. Abdoli M, Akbari M, Shahrabi J (2023) Bagging supervised autoencoder classifier for credit scoring. *Expert Syst Appl* 213:118991
115. Poutré C, Chételat D, Morales M (2024) Deep unsupervised anomaly detection in high-frequency markets. *J Finance Data Sci* 100129
116. Chen S, Guo W (2023) Auto-encoders in deep learning-a review with new perspectives. *Mathematics* 11(8):1777
117. Zhao Y, Chen J (2022) A survey on differential privacy for unstructured data content. *ACM Comput Surv (CSUR)* 54(10s):1–28
118. Prystavka P, Cholyshkina O, Dolgikh S, Karpenko D (2020) Automated object recognition system based on convolutional autoencoder. 2020 10th international conference on advanced computer information technologies. *IEEE*, pp 830–833
119. Li J, Xia C, Chen X (2017) A benchmark dataset and saliency-guided stacked autoencoders for video-based salient object detection. *IEEE Trans Image Process* 27(1):349–364
120. Joshi A, Fidalgo E, Alegre E, Fernández-Robles L (2019) Summocoder: an unsupervised framework for extractive text summarization based on deep auto-encoders. *Expert Syst Appl* 129:200–215
121. Shen T, Jia J, Li Y, Wang H, Chen B (2020) Enhancing music recommendation with social media content: an attentive multi-modal autoencoder approach. 2020 international joint conference on neural networks. *IEEE*, pp 1–8
122. Xu Z, Dang Y, Wang Q (2022) Potential buyer identification and purchase likelihood quantification by mining user-generated content on social media. *Expert Syst Appl* 187:115899
123. Lahasan B, Samma H (2022) Optimized deep autoencoder model for internet of things intruder detection. *IEEE Access* 10:8434–8448
124. Oussidi A, Elhassouny A (2018) Deep generative models: survey. 2018 International conference on intelligent systems and computer vision (ISCV). *IEEE*, pp 1–8
125. Nazabal A, Olmos PM, Ghahramani Z, Valera I (2020) Handling incomplete heterogeneous data using VAEs. *Pattern Recogn* 107:107501
126. Martinelli DD (2022) Generative machine learning for de novo drug discovery: a systematic review. *Comput Biol Med* 145:105403
127. Bandi A, Adapa PVSR, Kuchi YEVPK (2023) The power of generative AI: a review of requirements, models, input-output formats, evaluation metrics, and challenges. *Future Internet* 15(8):260
128. Ranjan A, Bolkart T, Sanyal S, Black MJ (2018) Generating 3D faces using convolutional mesh autoencoders. In: *Proceedings of the European conference on computer vision*, pp. 704–720
129. Wang Y, Dai B, Hua G, Aston J, Wipf D (2018) Recurrent variational autoencoders for learning nonlinear generative models in the presence of outliers. *IEEE J Select Topics Signal Process* 12(6):1615–1627
130. Zhao J, Kim Y, Zhang K, Rush A, LeCun Y (2018) Adversarially regularized autoencoders. *International conference on machine learning*. PMLR, pp 5902–5911
131. Gondara L (2016) Medical image denoising using convolutional denoising autoencoders. 2016 IEEE 16th international conference on data mining workshops. *IEEE*, pp 241–246

Publisher's Note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.